

Deploy a Continuous Machine Learning Model using Watson Studio.

In this lab, you will be using 2017 Chicago building data to make Chicago a safer place by building a model to predict when buildings are likely to fail inspection. You can then use the model to find which buildings are most dangerous and attend to those first. Continuous learning will be set up to monitor the model performance and adjust if necessary.

Goals:

In this lab you will

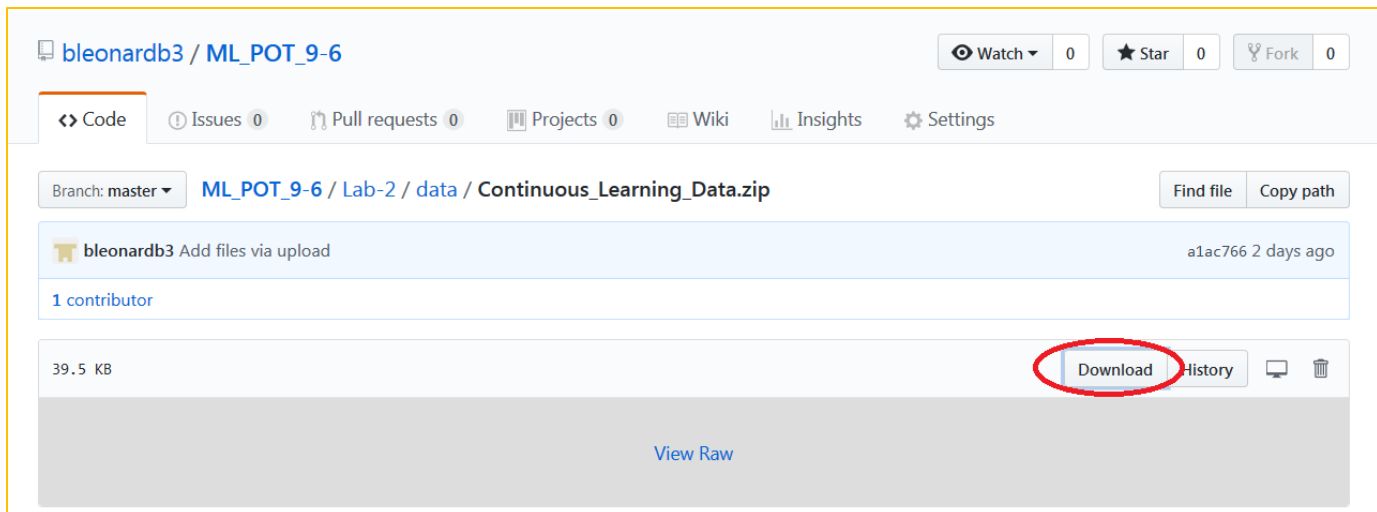
- Upload data to your Watson Studio project.
- Train, compare, and select a machine learning model.
- Set up continuous learning capabilities.
- Deploy a machine learning model.

Exercise Instructions

Step 1 Upload a New Data Asset.

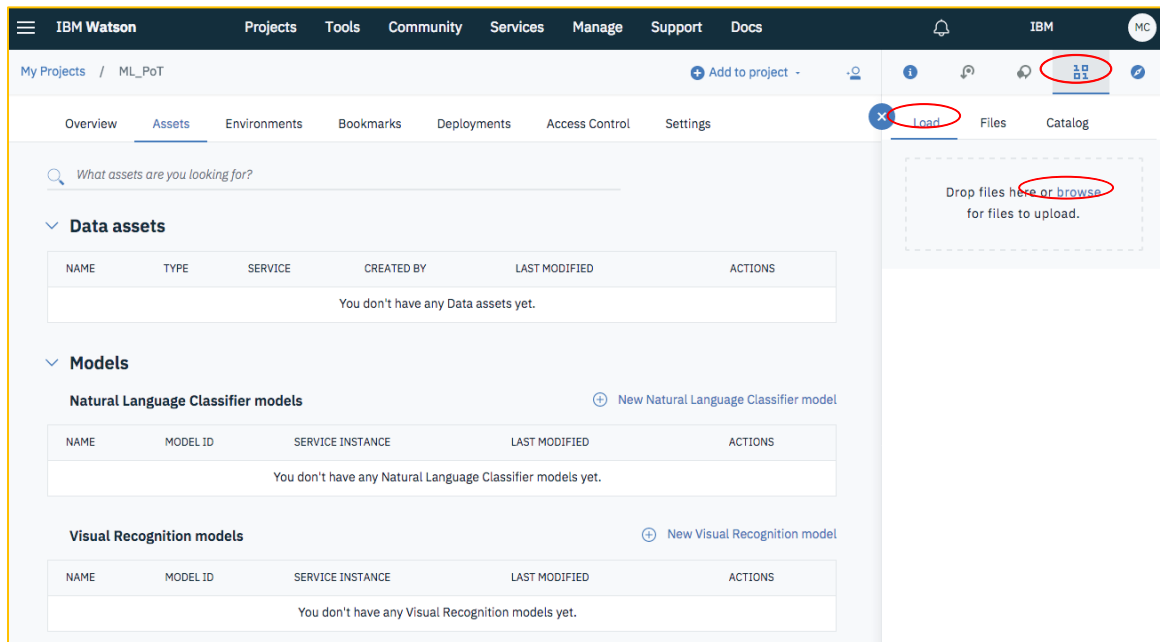
Before we build our models, we need to load data into our project.

1. Click on [Building Data](#)
2. Click on Download.



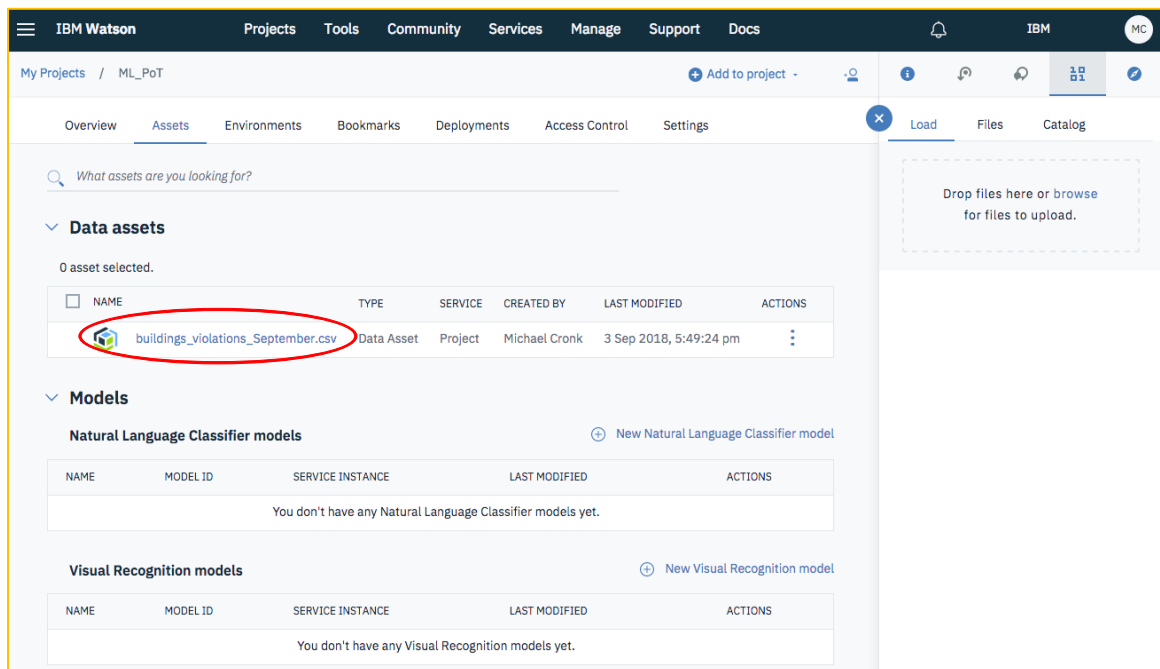
3. Save the zip file to your desktop and unzip the contents. This zip file contains data for building inspections for each month. We will begin by training our model for the month of September and then add October and November data to show how we can continuously adjust over time.
4. Return to your assets page. From Watson Studio, you can click on the Projects dropdown on the top navbar and then click on the name of your project.

5. In the top right of the screen, click the blue icon with 1's and 0's. This will open the data pallet.
6. Click on Load, browse, select the building_violations_September.csv file and click open.



The CSV file should now be listed under “Data assets.” It is now accessible by the Watson Studio modeling tools and applications.

7. Click on building_violations_September.csv under Data assets.



© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

After clicking on building_violations_September.csv, you are brought to a data view. Once you are satisfied that this is the data you want, click on your project name to return to your project's assets.

The screenshot shows the IBM Watson ML PoT interface. The navigation bar at the top includes 'My Projects', 'ML_PoT' (highlighted with a red circle), and 'buildings_violations_September.csv'. Below the navigation bar, there are tabs for 'Preview', 'Profile', and 'Lineage'. The 'Preview' tab is active, showing a table with 8 columns and 999 rows. The table columns are: INSPECTION_STAT..., VIOLATION_CO..., VIOLATION_STAT..., INSPECTION_CATEGO..., PROPERTY_GRO..., LONGITUDE, and LATITUDE. The table data shows various violation records. On the right side, there is a panel for the data asset 'buildings_violations_September.csv'. This panel includes a description, tags, and a 'Let's talk' button.

INSPECTION_STAT...	VIOLATION_CO...	VIOLATION_STAT...	INSPECTION_CATEGO...	PROPERTY_GRO...	LONGITUDE	LATITUDE
Type: String	Type: String	Type: String	Type: String	Type: String	Type: String	Type: String
FAILED	CN063014	OPEN	COMPLAINT	small	-87.691078	42.00245
FAILED	EL0086	OPEN	COMPLAINT	small	-87.714074	41.95654
FAILED	EL0086	OPEN	COMPLAINT	small	-87.633451	41.75931
FAILED	CN197087	OPEN	COMPLAINT	small	-87.795168	41.95265
FAILED	CN190019	OPEN	COMPLAINT	small	-87.769105	41.78514
FAILED	NC2022	OPEN	COMPLAINT	small	-87.712142	41.84204
FAILED	CN070014	OPEN	COMPLAINT	small	-87.713576	41.73967
FAILED	CN073024	COMPLIED	COMPLAINT	small	-87.727869	41.86917
FAILED	NC2022	OPEN	COMPLAINT	small	-87.830633	41.94983
FAILED	CN070034	OPEN	PERIODIC	small	-87.543687	41.73027
FAILED	PL157017	OPEN	COMPLAINT	small	-87.657112	41.79244
FAILED	PL237020	OPEN	COMPLAINT	small	-87.726443	41.90246
FAILED	VT1010	OPEN	PERMIT	small	-87.65453	41.92095
FAILED	EL0024	OPEN	COMPLAINT	small	-87.617871	41.80404
FAILED	PL194029	OPEN	COMPLAINT	small	-87.731726	41.92431
FAILED	CN197079	OPEN	COMPLAINT	small	-87.683636	41.88381
FAILED	CN141016	OPEN	COMPLAINT	small	-87.720044	41.94075

Schema: 8 Columns
Preview (999 rows)

Data Asset
buildings_violations_September.csv

Description
No description available for this asset

Tags
No tags available for this asset
Added: 09:49 PM UTC, 2018/09/03
Size: 71.679 KB

Let's talk

Step 2 Train, Compare, and Select a Machine Learning Model

1. Scroll down and click New Watson Machine Learning model.

Models

Natural Language Classifier models

New Natural Language Classifier model

NAME	MODEL ID	SERVICE INSTANCE	LAST MODIFIED	ACTIONS
You don't have any Natural Language Classifier models yet.				

Visual Recognition models

New Visual Recognition model

NAME	MODEL ID	SERVICE INSTANCE	LAST MODIFIED	ACTIONS
You don't have any Visual Recognition models yet.				

Watson Machine Learning models

New Watson Machine Learning model

NAME	STATUS	TYPE	RUNTIME	LAST MODIFIED		ACTIONS
You don't have any Watson Machine Learning models yet.						

Notebooks

New notebook

NAME	SHARED	SCHEDULED	STATUS	LANGUAGE	LAST EDITOR	LAST MODIFIED	ACTIONS
You don't have any Notebooks yet.							

© Copyright IBM Corp. 2018
Course materials may not be reproduced in whole or in part without the prior
written permission of IBM

2. Enter a Name and a Description. Note that the Machine Learning Service and the Spark Service should already be filled in from the services defined in the prerequisites. Select Manual, and click Create. We are creating a model using “Model builder” but can also create models in notebooks as well as with SPSS or Deep Learning flow modelers.

New model

Define model details

Name
Building_Violations_Chicago_2017

Description
Building violations continuous learning model.

Machine Learning Service
pm-20-pw

Select model type

☒ Model builder ☐ From file ☐ From sample

Spark Service or Environment
Only Spark environments supporting Scala kernels can be used for model builder creation.
spark-nn

Automatic
Prepare my data and create a model automatically

Manual
Let me prepare my data and select which models to train

Need something more flexible? Create a [notebook](#) or design a [Modeler flow](#)

Let's talk

Cancel **Create**

3. You will be brought to a “Select data asset” page. Select building_violations_September.csv as our data asset and click Next.

Select data asset

The model builder currently supports CSV files and IBM Db2 Warehouse on Cloud data assets.

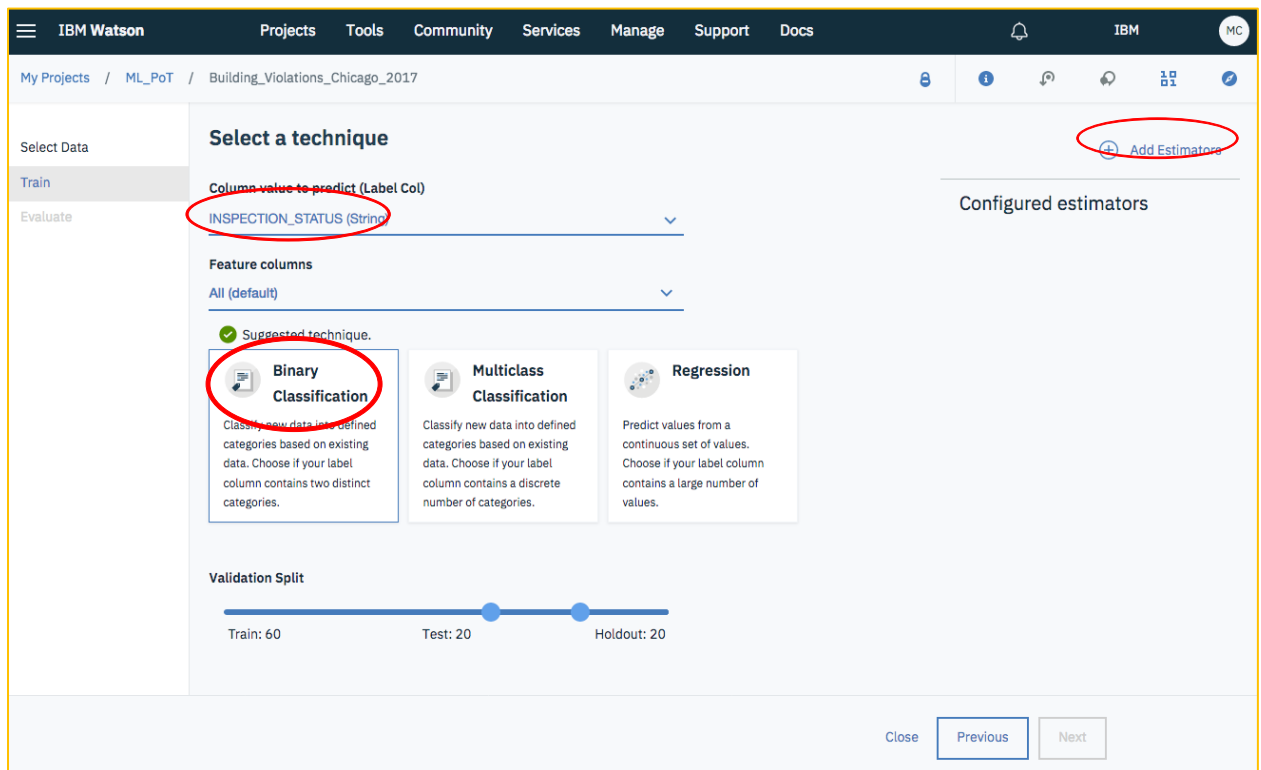
What asset are you looking for?

NAME	TYPE	SERVICE
<input checked="" type="radio"/> buildings_violations_September.csv	Data Asset	Project

Close **Next**

We will now choose which techniques may work well given our data and given our goal to predict which buildings are most likely to not pass inspection.

4. On the “Select a technique” page, click on the Select Label Col dropdown and select `INSPECTION_STATUS`(String).
5. Select Binary Classification. Note that this is often automatically suggested by Watson Studio.
6. Click Add Estimators in the upper right corner of the page.

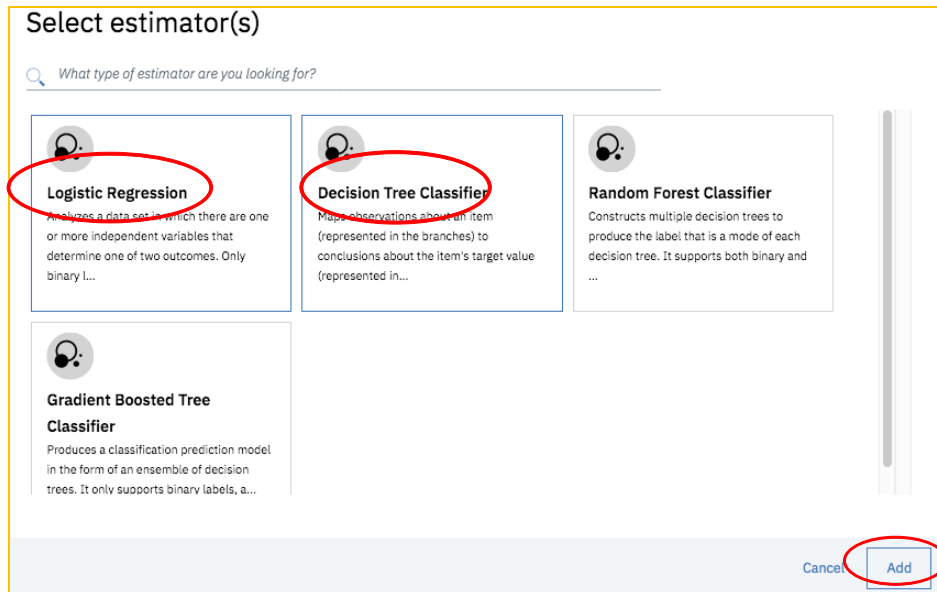


© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

Watson Studio only presents the estimators that would work with our selections on the previous page.

7. Select Logistic Regression and Decision Tree Classifier.
8. Click Add.



9. Note that you can click on the configured estimators to tune the hyperparameters associated with the algorithm. Click on Logistic Regression.

Select Data

Train

Evaluate

Select a technique

You cannot change label column, feature columns, model type, or validation split after adding an estimator.
You must first delete all estimators in order to make changes to these attributes.

Column value to predict (Label Col)
INSPECTION_STATUS (String)

Feature columns
All (default)

Configured estimators

Logistic Regression
Not Yet Trained

Decision Tree Classifier
Not Yet Trained

Add Estimator

10. We will not change any hyperparameters, so click on Cancel.

Configure Logistic Regression

Weight column
Select

Elastic net parameter (double)
0

Fit an intercept
☒ fitIntercept

Maximum iterations for convergence (integer)
100

Regularization parameter (double)
0

Standardize the training features before fitting the model
☒ standardization

Threshold (double)
0.5

Cancel Save

11. Back on the “Select a technique” page, click Next.

Select a technique

You cannot change label column, feature columns, model type, or validation split after adding an estimator. You must first delete all estimators in order to make changes to these attributes.

Column value to predict (Label Col)
INSPECTION_STATUS (String)

Feature columns
All (default)

Suggested technique.

Binary Classification

Classify new data into defined categories based on existing data. Choose if your label column contains two distinct categories.

Multiclass Classification

Classify new data into defined categories based on existing data. Choose if your label column contains a discrete number of categories.

Regression

Predict values from a continuous set of values. Choose if your label column contains a large number of values.

Validation Split

Train: 60 Test: 20 Holdout: 20

Configured estimators

- Logistic Regression
Not Yet Trained
- Decision Tree Classifier
Not Yet Trained

Close Previous **Next** Let's talk

The Select model page will allow us to compare the results of different estimator types.

12. Both models have performed well. For this tutorial, select Logistic Regression.


13. Click Save.

Select model

ESTIMATOR TYPE	STATUS	PERFORMANCE	AREA UNDER ROC CURVE	AREA UNDER PR CURVE	LAST EVALUATION	ACTIONS
<input checked="" type="radio"/> LogisticRegression	Trained & Evaluated	Excellent	0.95123	0.61348	3 Sep 2018, 6:00 PM	⋮
<input type="radio"/> DecisionTreeClassifier	Trained & Evaluated	Excellent	0.90136	0.53478	3 Sep 2018, 6:00 PM	⋮

Close Previous **Save**

14. When the “Save model” window appears, click Save.



We have just saved the model in our Watson Machine Learning service and can now view information about the model specifications, details of construction, and input schema. The model can also be exposed as an API and used by data pipelines, applications, or other external systems.

15. Click on the console button with an “angle bracket” to view the schema in JSON format.

The screenshot shows the IBM Watson Machine Learning console interface. The top navigation bar includes "IBM Watson", "Projects", "Tools", "Community", "Services", "Manage", "Support", and "Docs". The breadcrumb trail is "My Projects / ML_PoT / Building_Violations_Chicago_2017". The main heading is "Building_Violations_Chicago_2017" with a trash icon. Below this are tabs for "Overview", "Evaluation", "Deployments", and "Lineage". The "Overview" tab is active, showing a "Summary" section with a table of model details. Below the summary is the "Input Schema" section, which contains a table of columns and their types. In the top right corner of the "Input Schema" table, there is a button with a table icon and an angle bracket icon, which is circled in red. A "Let's talk" button is located at the bottom right of the console.

COLUMN	TYPE
VIOLATION_CODE	string
VIOLATION_STATUS	string
INSPECTION_CATEGORY	string
PROPERTY_GROUP	string

© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

This schema can be copied and used elsewhere to help existing systems easily interact with our model and the Watson Machine Learning service.

Input Schema

```
{
  {
    "name": "VIOLATION_CODE",
    "type": "string",
    "nullable": true,
    "metadata": {
      "columnInfo": {
        "columnPrimaryKey": false,
        "columnTypeName": "varchar",
        "columnSigned": true,
        "columnType": 12,
        "columnLength": 1024,
        "columnNullable": true,
        "columnScale": 0
      }
    }
  }
}
```

16. Scroll up and click on the Evaluation tab.

IBM Watson Projects Tools Community Services Manage Support Docs

My Projects / ML_PoT / Building_Violations_Chicago_2017

Building_Violations_Chicago_2017

Overview **Evaluation** Deployments Lineage

Summary

Machine learning service	pm-20-pw
Model Type	wml-1.1
Runtime environment	spark-2.1
Training date	3 Sep 2018, 5:03 PM
Label column	INSPECTION_STATUS
Latest version	2456f175-bee0-421a-901a-b414d4b61534

Input Schema


COLUMN	TYPE
VIOLATION_CODE	string
VIOLATION_STATUS	string
INSPECTION_CATEGORY	string
PROPERTY_GROUP	string

Let's talk

Step 3 Set up Continuous Learning Model Capabilities

Model performance is critical for solving data science problems. This page provides information and functionality to aid in continuous training, tuning, and redeployment. We will now set up parameters to automatically retrain our model when performance falls below a certain threshold.

1. Click on Configure Performance Monitoring.



The screenshot shows the IBM Watson ML interface for a project named 'Building_Violations_Chicago_2017'. The 'Evaluation' tab is selected, showing the 'Last Evaluation Result' table and the 'Performance Monitoring' section. The 'Configure Performance Monitoring' link is circled in red.

Version	2456f175-bee0-421a-901a-b414d4b61534
Phase	setup
AreaUnderPR	0.902
AreaUnderROC	0.989

Performance Monitoring

Configure performance monitoring to evaluate and retrain the model periodically to ensure the model performance is acceptable. You will need an existing IBM Db2 Warehouse on Cloud connection associated with your project to be used as your feedback data connection.

[Configure Performance Monitoring](#)

Versions

TIME	VERSION	DEPLOYED	ACTIONS
03 Sep 2018 05:03pm	2456f175-bee0-421a-901a-b414d4b61534		

Let's talk

2. Under Prediction type, select binary.
3. Under Metric details, select areaUnderPR and enter 0.8.

This relies on a database table of new data. We will use DB2 Warehouse on Cloud, a data store optimized for analytic data sets.

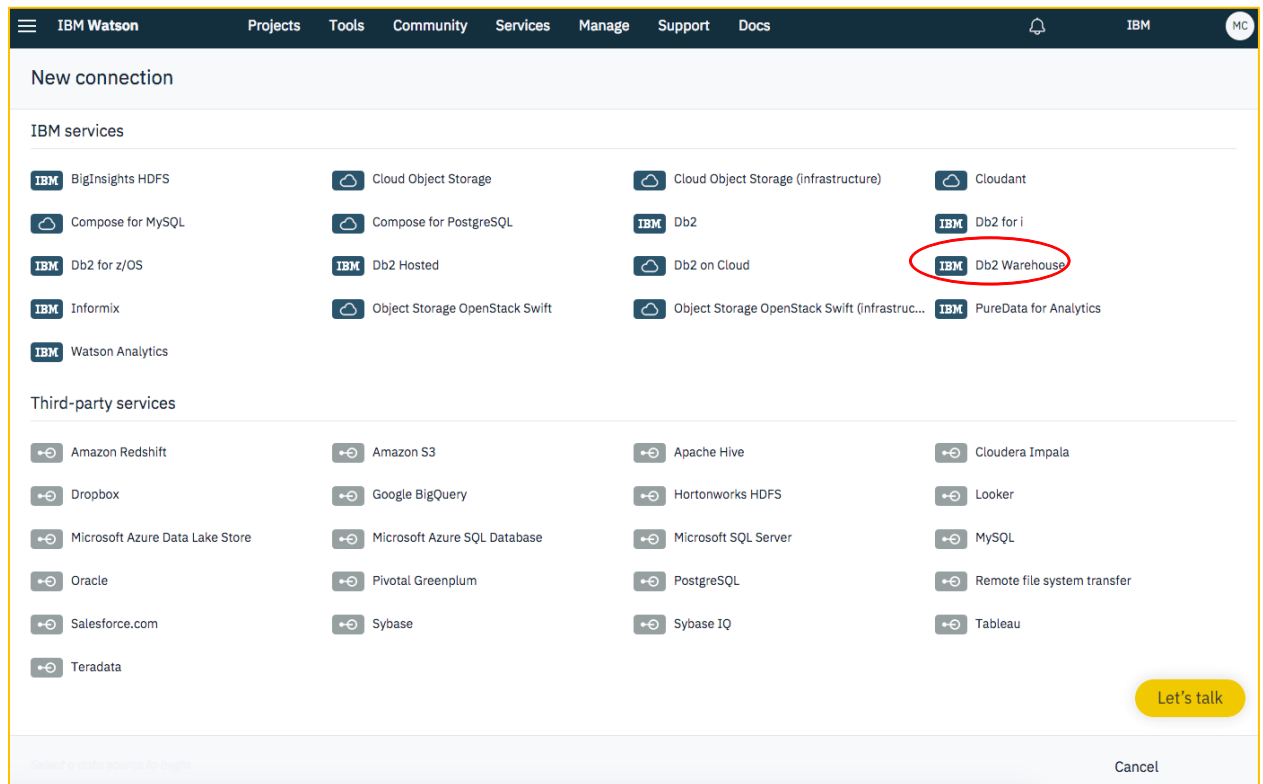
4. Click on Create a new connection.

The screenshot shows the 'Configure performance monitoring' page in the IBM Watson interface. The page has a dark blue header with navigation links: IBM Watson, Projects, Tools, Community, Services, Manage, Support, Docs, and a user profile icon. The main content area is white and contains several sections:

- Spark Service or Environment:** A dropdown menu showing 'spark-nn'.
- Prediction type:** A dropdown menu showing 'binary'.
- Metric details (type / optional threshold):** A dropdown menu showing 'areaUnderPR' and a text input field showing '0.8'.
- Feedback data connection (IBM Db2 Warehouse on Cloud):** A dropdown menu showing 'Create new connection'.
- Record count required for re-evaluation:** A text input field showing '1000'.
- Auto retrain:** A dropdown menu showing 'when model performance is below threshold'.
- Auto deploy:** A dropdown menu showing 'when model performance is better than previous version'.

At the bottom right of the page, there are three buttons: 'Cancel', 'Save', and 'Let's talk'.

5. On the “New connection” page, click DB2 Warehouse.



6. Access [DB2 Credentials](#)

Ask your instructor for your assigned database.

7. Enter a Name for your connection (e.g. DB2 Warehouse Connection).

8. Using the credentials cut and paste the fields under Hostname, Database, Username, and Password.

The screenshot shows the 'Connection overview' and 'Connection details' form. The 'Name' field is circled in red and contains 'Db2 Warehouse Connection'. The 'Description' field contains 'IBM Db2 warehouse database on Cloud'. The 'Connection details' section includes fields for 'Database *', 'Password *', 'Hostname or IP Address *', and 'Username *'. The 'Database' field is circled in red and contains 'RLUDB'. The 'Password' field is circled in red and contains '*****'. The 'Hostname or IP Address' field is circled in red and contains 'dashdb-entry-yp-dal09-08.services.dal.bj'. The 'Username' field is circled in red and contains 'dash11054'. The 'Secure Gateway' checkbox is unchecked.

© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

9. Click Create.

New connection (Db2 Warehouse Connection - Db2 Warehouse)

Connection overview

Name

Db2 Warehouse Connection

Description

IBM Db2 warehouse database on Cloud

2965

Connection details

Database *

BLUDB

Password *

.....

Secure Gateway

☐ Use a secure gateway

Hostname or IP Address *

dashdb-entry-yp-dal09-08.services.dal.bl

Username *

dash11054

Let's talk

Cancel

Create

10. Return to the “Configure performance monitoring page” and click Select feedback data reference.

The screenshot shows the 'Configure performance monitoring' page in the IBM Watson interface. The page has a dark blue header with the IBM Watson logo and navigation links: Projects, Tools, Community, Services, Manage, Support, Docs. There is a notification bell icon, the IBM logo, and a user profile icon labeled 'MC'.

The main content area is titled 'Configure performance monitoring'. It contains several sections:

- Spark Service or Environment**: A dropdown menu showing 'spark-nn'. Below it, a note states: 'Only Spark environments supporting Scala kernels can be used for continuous learning.'
- Prediction type**: A dropdown menu showing 'binary'.
- Metric details (type / optional threshold)**: A dropdown menu showing 'areaUnderPR' and a text input field showing '0.8' with a plus icon to its right.
- Feedback data connection (IBM Db2 Warehouse on Cloud - Create new connection [3])**: This section is highlighted with a red circle. It contains a link 'Select feedback data reference'.
- Record count required for re-evaluation**: A text input field showing '1000'.
- Auto retrain**: A dropdown menu showing 'when model performance is below threshold'.
- Auto deploy**: A dropdown menu showing 'when model performance is better than previous version'.

At the bottom right of the page, there are three buttons: 'Cancel', 'Save', and 'Let's talk' (which is highlighted in yellow).

11. Click on your DB2 Warehouse-xx name.
12. Select the schema that matches DASH####.
13. Click select.

Select feedback data reference

ML_PoT	Db2 Warehouse Connection	DASH11054
Connections (1)	Schemas (13)	
Db2 Warehouse Connection	DASH11054	
	ERRORSCHEMA	No drilldowns currently exist.
	GOSALES	
	GOSALESBW	
	GOSALESHR	
	GOSALESMR	
	GOSALESRT	
	IBM_RTMON_DATA	
	NULLIDR1	
	NULLIDRA	
	SAMPLES	
	ST_INFORMTN_SCHEMA	
	TEST	

Cancel Select Let's talk

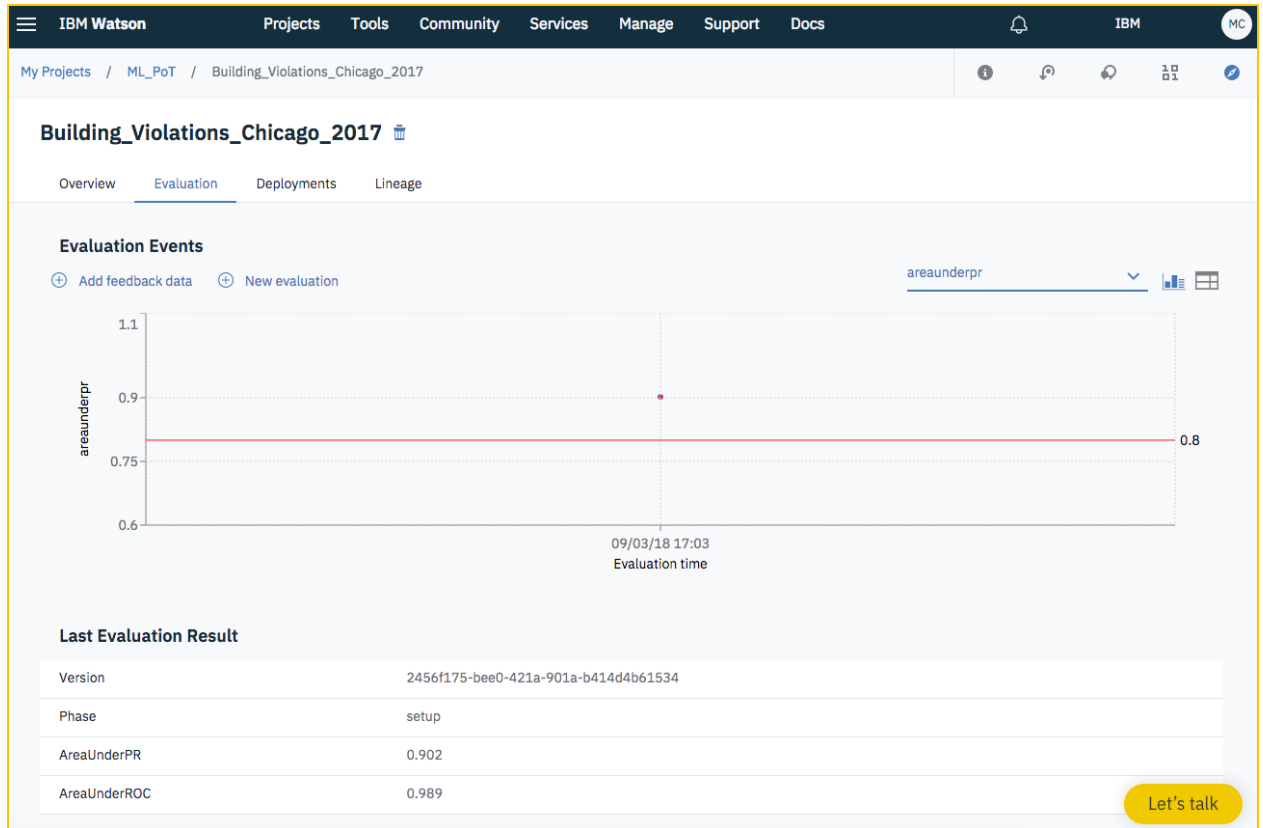
14. After returning to the “Configure performance monitoring” page, enter a unique table name (attendees are sharing the DB2 Warehouse service so make your name unique by appending your initials to the table name).
15. Enter 500 as the Record count.
16. Under Auto deploy, select never.
17. Click Save.

The screenshot shows the 'Configure performance monitoring' page in the IBM Watson interface. The page has a dark blue header with navigation links: IBM Watson, Projects, Tools, Community, Services, Manage, Support, Docs, a notification bell, IBM, and a user profile icon (MC). The main content area is white and contains the following sections:

- Spark Service or Environment**: Only Spark environments supporting Scala kernels can be used for continuous learning. The dropdown menu shows 'spark-nn'.
- Prediction type**: The dropdown menu shows 'binary'.
- Metric details (type / optional threshold)**: The dropdown menu shows 'areaUnderPR' and the threshold is '0.8'.
- Feedback data connection**: (IBM Db2 Warehouse on Cloud - Create new connection [\[+\]](#))
dashdb: BLUDB [Change feedback data reference](#)
New2017Table
- Record count required for re-evaluation**: 500
- Auto retrain**: when model performance is below threshold
- Auto deploy**: never

At the bottom right, there are three buttons: 'Cancel', 'Save', and 'Let's talk'. The 'Save' button is highlighted with a red circle.

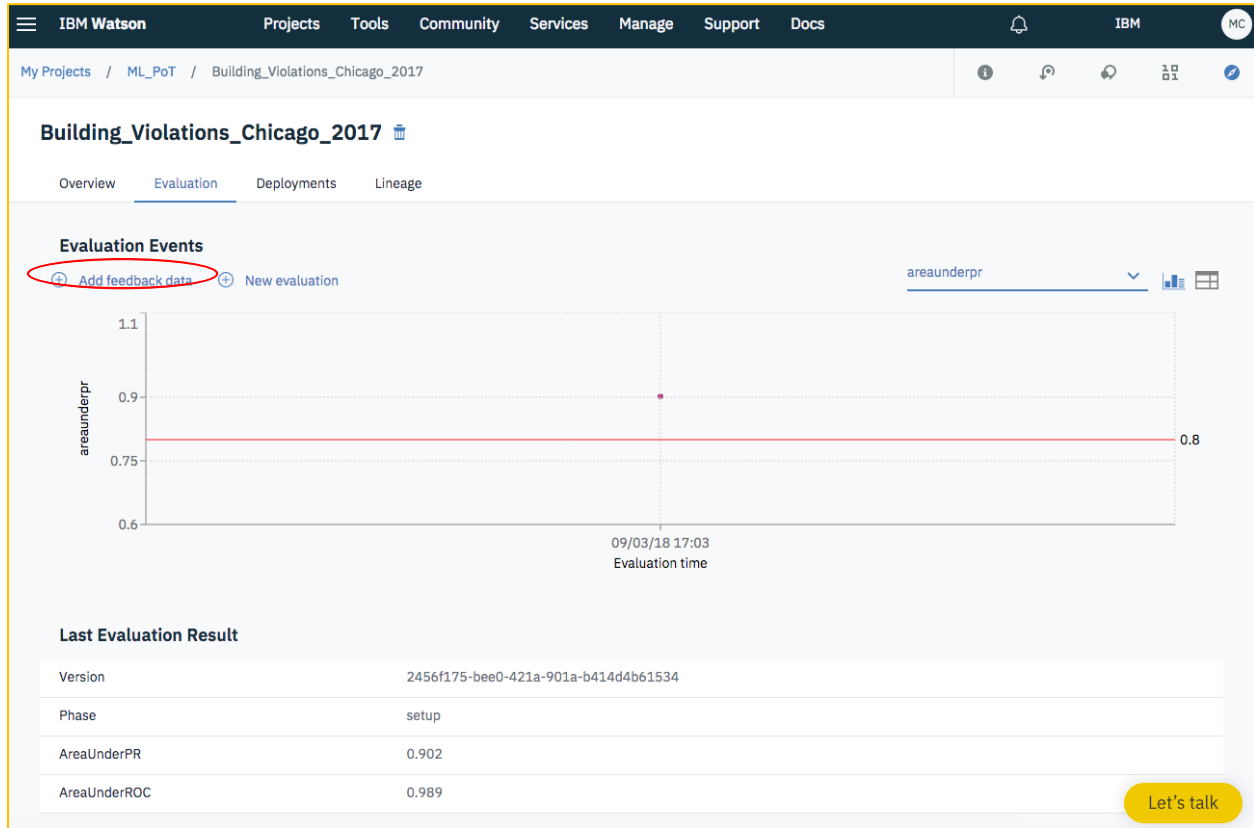
We have just set our model to retrain whenever its performance falls below 0.8. The evaluation screen has changed to reflect the decisions we have just made. Notice the graph lists our current areaunderPR and the threshold we have set. Clicking on Add feedback data will upload a new dataset. Clicking on new evaluation will then evaluate the model and check to see how the model compares to our 0.8 threshold.



Suppose September has passed, and we now have October data. We can trigger a new evaluation by adding more feedback data.

18. Click Add feedback data.

19. Select building_violations_October.csv.




© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

20. When the “New evaluation” window appears, click New Evaluation. Re-evaluating the model may take a few minutes. This is a great time to grab coffee.

Notice that the October data has caused our model to fall below the 0.8 threshold and a new version has been trained. Although we run through these steps graphically, you can also configure a Watson Machine Learning continuous learning model entirely through APIs.



You can continue by adding the November data by the same process of hitting **Add feedback data** and adding more files. If you scroll down you will be able to see all model versions listed. We have three versions, but no deployments.

My Projects / ML_PoT / Building_Violations_Chicago_2017

Last Evaluation Result

Version	028ee821-4695-462f-9ec8-6d85dee03cee
Phase	training
AreaUnderPR	0.726

Performance Monitoring [Edit configuration](#)

Performance Metrics (Threshold)	areaUnderPR (0.8)
Feedback Data Reference	dashdb: BLUDB / New2017Table
Training Data Reference	COS: mlpot-donotdelete-pr-d2mpgp7cdyrjlt
Record Count Required For Re-evaluation	500
Auto Re-train	conditionally
Auto Re-deploy	never

Versions

TIME	VERSION	DEPLOYED	AREAUNDERPR	ACTIONS
03 Sep 2018 05:33pm	028ee821-4695-462f-9ec8-6d85dee03cee		0.726	...
03 Sep 2018 05:27pm	15417ea0-2fe4-4b74-b053-53dbda504256		0.682	...
03 Sep 2018 05:03pm	2456f175-bee0-421a-901a-b414d4b61534		0.902	...

Let's talk


© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

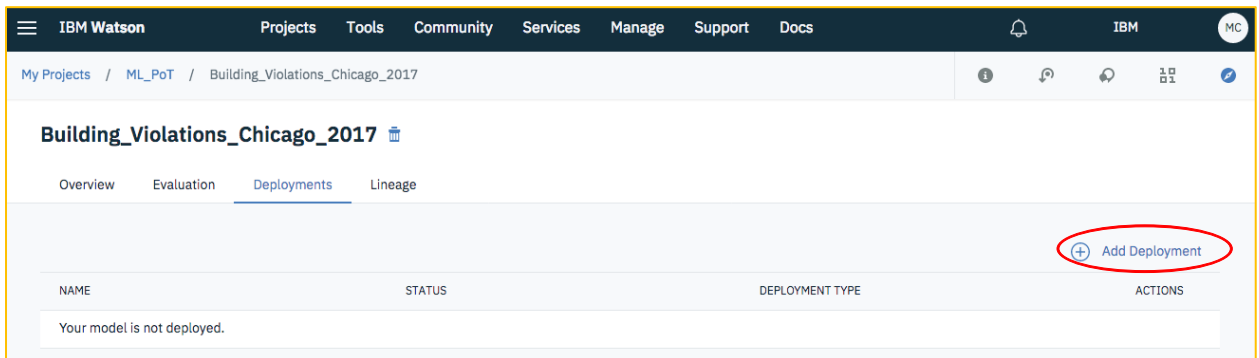
Step 4 Deploy a Machine Learning Model.

Now we will focus on deployment.

1. Click the Deployment tab.



2. Add Deployment.



3. Enter a Name (e.g. ChicagoOnline).
4. Enter a Description.
5. Under Deployment type, select Web service.
6. Click Save.

IBM Watson Projects Tools Community Services Manage Support Docs

Create Deployment

Define deployment details

Name
ChicagoOnline

Description
Building violations 2017 model deployment

Deployment type
☒ Web service
☐ Batch prediction
☐ Realtime streaming prediction

Cancel Save Let's talk

The model is now accessible by external systems.

7. Click on your deployment (e.g. ChicagoOnline).

IBM Watson Projects Tools Community Services Manage Support Docs

My Projects / ML_PoT / Building_Violations_Chicago_2017

Building_Violations_Chicago_2017

Overview Evaluation Deployments Lineage

+ Add Deployment

NAME	STATUS	DEPLOYMENT TYPE	ACTIONS
ChicagoOnline	DEPLOY_SUCCESS	Web Service	

© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

You are automatically brought to the Overview tab lists information related to the model deployment including services used and version number.

8. Click the Implementation tab.

The screenshot shows the IBM Watson ML console interface. The top navigation bar includes 'IBM Watson', 'Projects', 'Tools', 'Community', 'Services', 'Manage', 'Support', 'Docs', a notification bell, 'IBM', and a user profile icon 'MC'. The breadcrumb trail is 'My Projects / ML_PoT / Building_Violations_Chicago_2017 / ChicagoOnline'. The main content area is titled 'ChicagoOnline' and has three tabs: 'Overview', 'Implementation' (highlighted with a red circle), and 'Test'. Below the tabs, there are two sections: 'Deployment' and 'Model'. The 'Deployment' section contains a table with the following data:

Name	ChicagoOnline
Type	Web Service
Deployment ID	5d232eda-13b3-440f-8e13-de9ee1ea2187
Status	DEPLOY_SUCCESS
Asset type	model
Asset name	Building_Violations_Chicago_2017
Machine learning service	pm-20-pw
Created	03 Sep 2018 05:38pm
Last modified	03 Sep 2018 05:39pm

The 'Model' section contains a table with the following data:

Name	Building_Violations_Chicago_2017
ID	ffa1749a-c5e3-49e6-91fe-212b16e97dfb
Version ID	2456f175-bee0-421a-901a-b414d4b61534

A yellow 'Let's talk' button is located in the bottom right corner of the console area.

The implementation tab provides developers information to help minimize the time it takes to develop models and place them in a production environment.

9. Click on the Test tab.

The screenshot displays the IBM Watson ChicagoOnline web interface. The top navigation bar includes 'IBM Watson', 'Projects', 'Tools', 'Community', 'Services', 'Manage', 'Support', and 'Docs'. Below this, a breadcrumb trail shows 'My Projects / ML_PoT / Building_Violations_Chicago_2017 / ChicagoOnline'. The main content area is titled 'ChicagoOnline' and features three tabs: 'Overview', 'Implementation', and 'Test'. The 'Test' tab is highlighted with a red circle. The 'Implementation' section contains a table with the following details:

Scoring End-point	https://us-south.ml.cloud.ibm.com/v3/wml_instances/2b0bce0d-7638-4b40-a2b2-0d4ed148368a/deployments/5d232eda-13b3-440f-8e13-de9ee1ea2187/online
Authorization: Bearer <token>	See code snippets below for information on how to retrieve the WML Authorization Token to be passed with scoring requests.
Content-type: application/json	Required if the request body is sent in JSON format.

Below the table, there is a 'Code Snippets' section with tabs for 'cURL', 'Java', 'JavaScript', 'Python', and 'Scala'. The 'cURL' tab is selected, showing the following code:

```
# retrieve your $WML_SERVICE_CREDENTIALS_USERNAME, $WML_SERVICE_CREDENTIALS_PASSWORD, and $WML_SERVICE_CREDENTIALS_URL from the
# Service credentials associated with your IBM Cloud Watson Machine Learning Service instance.

curl --basic --user $WML_SERVICE_CREDENTIALS_USERNAME:$WML_SERVICE_CREDENTIALS_PASSWORD $WML_SERVICE_CREDENTIALS_URL/v3/identity/token

# the above CURL request will return an auth token that you will use as $WML_AUTH_TOKEN in the scoring request below
# TODO: manually define and pass values to be scored below
curl -X POST --header 'Content-Type: application/json' --header 'Accept: application/json' --header "Authorization: Bearer $WML_AUTH_TOKEN" -d '{
  "text": "The cat sat on the mat."
}'
```

The Test tab allows manual testing of the deployed model and viewing of results.

10. Enter:

VIOLATION_CODE: CN063014

VIOLATION_STATUS: OPEN

INSPECTION_CATEGORY: COMPLAINT

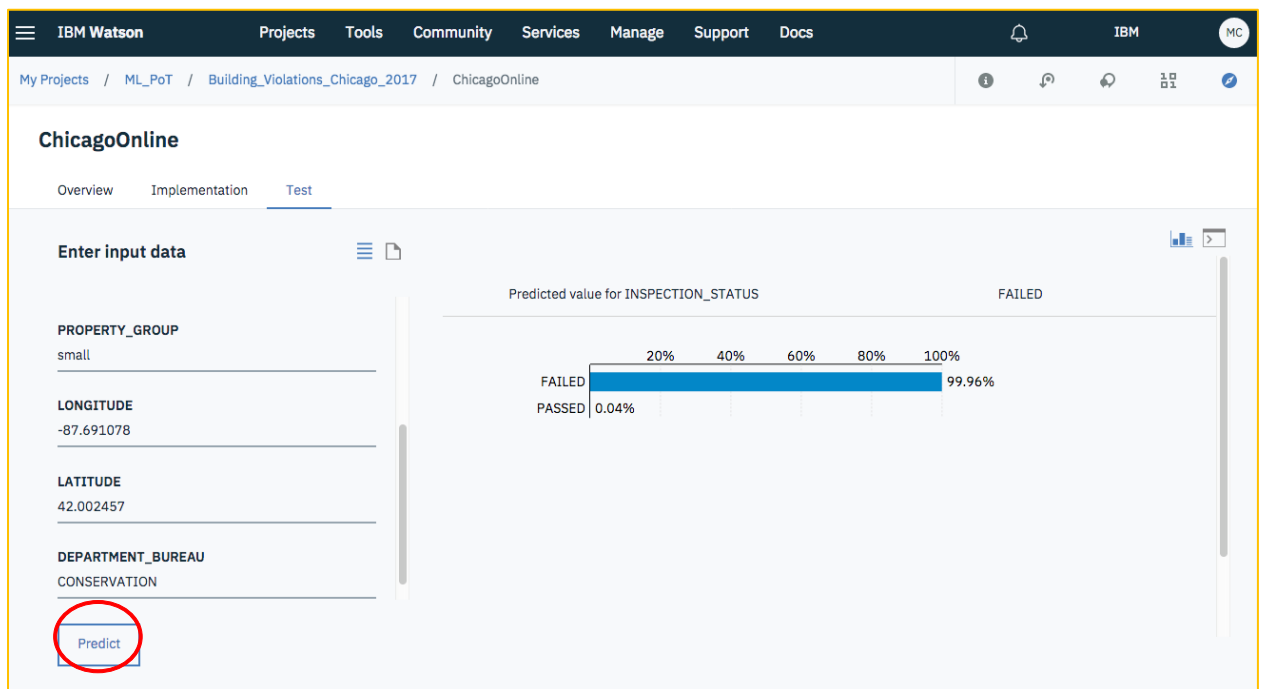
PROPERTY_GROUP: small

LONGITUDE: -87.691078

LATITUDE: 42.002457

DEPARTMENT_BUREAU: CONSERVATION

11. Click Predict.




© Copyright IBM Corp. 2018

Course materials may not be reproduced in whole or in part without the prior written permission of IBM

According to our model, a building with our input data has a 0.23% chance of passing inspection.

12. Feel free to change the input data and run new predictions. When ready, click on your model name (e.g. Building_Violations_Chicago_2017)



13. Click on the Evaluation tab.

The screenshot shows the IBM Watson Machine Learning interface. The top navigation bar includes 'IBM Watson', 'Projects', 'Tools', 'Community', 'Services', 'Manage', 'Support', and 'Docs'. The breadcrumb trail is 'My Projects / ML_PoT / Building_Violations_Chicago_2017'. The project name 'Building_Violations_Chicago_2017' is displayed with a trash icon. Below the project name, there are four tabs: 'Overview', 'Evaluation' (which is circled in red), 'Deployments', and 'Lineage'. The 'Evaluation' tab is active, showing a 'Summary' section with a table of model details.

Summary	
Machine learning service	pm-20-pw
Model Type	wml-1.1
Runtime environment	spark-2.1
Training date	3 Sep 2018, 5:03 PM
Label column	INSPECTION_STATUS
Latest version	028ee821-4695-462f-9ec8-6d85dee03cee

If you scroll to the bottom you will be able to see the different model versions that have been created and which model is currently deployed.

Versions				
TIME	VERSION	DEPLOYED	AREAUNDERPR	ACTIONS
03 Sep 2018 05:33pm	028ee821-4695-462f-9ec8-6d85dee03cee		0.726	⋮
03 Sep 2018 05:27pm	15417ea0-2fe4-4b74-b053-53bda504256		0.682	⋮
03 Sep 2018 05:03pm	2456f175-bee0-421a-901a-b414d4b61534	✓	0.902	⋮

Let's talk

During this lab we have very quickly compared various machine learning models and chosen the best one tuned to our dataset and objectives. We then created a continuous machine learning model that automatically monitors and retrains allowing Watson Machine Learning to keep applications, data pipelines, or external systems relying on the machine learning model as up to date as possible.

You have completed Lab 02